

Poisoning_Attack_Spam_Filter

April 22, 2022

1 Poisoning Attack Spam Filter (Spam or Ham)

Code based on [Build a machine learning email spam detector with Python](<https://blog.logrocket.com/email-spam-detector-python-machine-learning/>)

1.1 Import Libraries

```
[2]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.feature_extraction.text import CountVectorizer
      from sklearn import svm
      import matplotlib.pyplot as plt
      from sklearn.metrics import ConfusionMatrixDisplay
      import random
```

1.2 Download Data and Process CSV

```
[3]: !wget https://raw.githubusercontent.com/SmallLion/Python-Projects/main/
      ↪Spam-detection/spam.csv
```

```
--2022-04-22 14:04:39-- https://raw.githubusercontent.com/SmallLion/Python-
Projects/main/Spam-detection/spam.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.111.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com
(raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 499362 (488K) [text/plain]
Saving to: 'spam.csv'
```

```
spam.csv          100%[=====>] 487.66K  --.-KB/s    in 0.04s
```

```
2022-04-22 14:04:39 (12.9 MB/s) - 'spam.csv' saved [499362/499362]
```

```
[4]: spam = pd.read_csv('spam.csv')
```

1.3 Split Dataset into Train and Test

```
[14]: print("Number of 'HAM' instances:", len([ y for y in spam["v1"] if y == 'ham']))  
      ↪  
      print("Number of 'SPAM' instances:", len([ y for y in spam["v1"] if y ==  
      ↪ 'spam'])) )
```

Number of 'HAM' instances: 4825

Number of 'SPAM' instances: 747

```
[5]: x = spam['v2']  
     y = spam["v1"]  
     x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2)  
     x_train, x_poisoning, y_train, y_poisoning = train_test_split(x_train,  
     ↪ y_train, test_size = 0.5)
```

```
[6]: data_train_dic = {'x': x_train, 'y': y_train}  
     data_train = pd.DataFrame(data_train_dic)  
     data_test_dic = {'x': x_test, 'y': y_test}  
     data_test = pd.DataFrame(data_test_dic)  
     data_poisoning_dic = {'x': x_poisoning, 'y': y_poisoning}  
     data_poisoning = pd.DataFrame(data_poisoning_dic)
```

```
[7]: print("TRAIN DATA. size:", len(data_train))  
     print(data_train.head(10))
```

TRAIN DATA. size: 2228

		x	y
4856	Thanks and ! Or bomb and date as my phone wan...		ham
954	Filthy stories and GIRLS waiting for your		spam
1737	I cant pick the phone right now. Pls send a me...		ham
2571	%Ū_ and don%Ū÷t worry we%Ū÷ll have finished by...		ham
1109	S s..first time..dhoni rocks...		ham
2752	Sat right? Okay thanks...		ham
4270	Natalja (25/F) is inviting you to be her frien...		spam
2990	HOT LIVE FANTASIES call now 08707509020 Just 2...		spam
4414	Yeah I should be able to, I'll text you when I...		ham
4755	Don't make life too stressfull.. Always find t...		ham

```
[8]: print("TEST DATA. size:", len(data_test))  
     print(data_test.head(10))
```

TEST DATA. size: 1115

		x	y
3204	Oh, yes, I've just been a little under the wea...		ham
3604	Jordan got voted out last nite!		ham
1986	S..antha num corrcr dane		ham
3406	Beautiful Truth against Gravity.. Read careful...		ham

```

706                So when do you wanna gym harri ham
3662 Frnd s not juz a word...not merely a relatio... ham
5244 thanks for the temales it was wonderful. Thank... ham
2820                Oh god..taken the teeth?is it paining ham
4811        fyi I'm at usf now, swing by the room whenever ham
4869 Dip's cell dead. So i m coming with him. U bet... ham

```

```

[9]: print("POISONING DATA. size:", len(data_poisoning))
      print(data_poisoning.head(10))

```

POISONING DATA. size: 2229

```

                                     x      y
3361        No messages on her phone. I'm holding it now ham
622  MAKE SURE ALEX KNOWS HIS BIRTHDAY IS OVER IN F... ham
203                                     Goodmorning sleeping ga. ham
1889                I gotta collect da car at 6 lei. ham
423  URGENT! Your Mobile number has been awarded wi... spam
2493 they released another Italian one today and it... ham
999                Aight will do, thanks again for comin out ham
2814 Some are lasting as much as 2 hours. You might... ham
5048 Edison has rightly said, \A fool can ask more ... ham
1490 Cant believe i said so many things to you this... ham

```

1.4 Train Model

```

[10]: cv = CountVectorizer()
      features = cv.fit_transform(x_train)

```

Support Vector Machines for Classification: finds a function that best separates both categories.

```

[11]: model = svm.SVC()
      model.fit(features,y_train)

```

```

[11]: SVC()

```

1.5 Evaluate Model Accuracy

```

[15]: features_test = cv.transform(x_test)
      print("Accuracy: {}".format(model.score(features_test,y_test)))

```

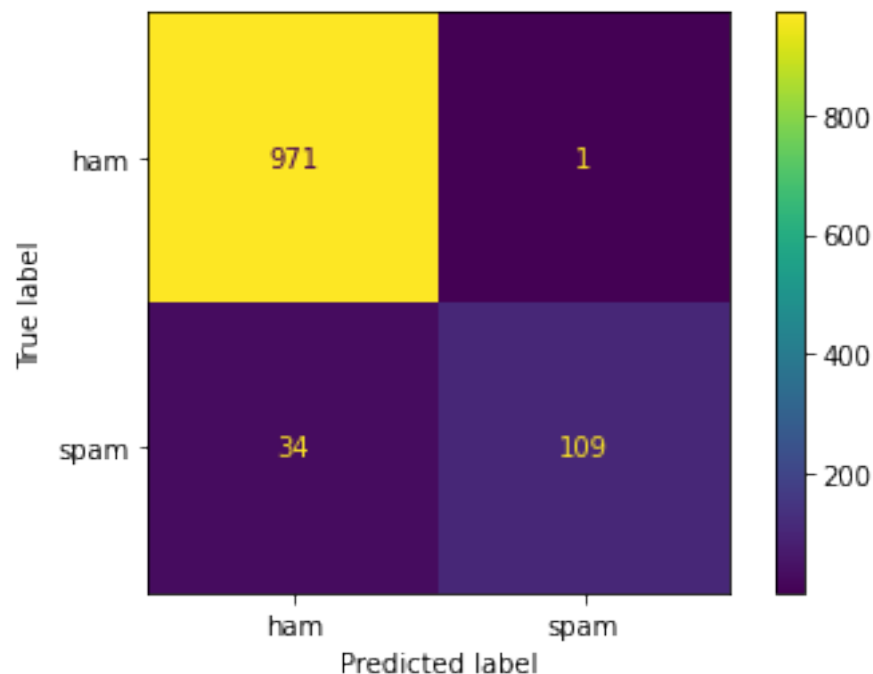
Accuracy: 0.968609865470852

```

[16]: y_pred_test = model.predict(features_test)

      ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
      plt.show()

```



1.6 Poisoning of Training Data and Re-train model

```
[17]: for i, y in enumerate(y_poisoning):
      if y == 'ham':
          y_poisoning.iloc[i] = 'spam'
      else:
          y_poisoning.iloc[i] = 'ham'
```

```
[18]: data_poisoning_dic = {'x': x_poisoning, 'y': y_poisoning}
      data_poisoning = pd.DataFrame(data_poisoning_dic)
      print("POISONED DATA. size:", len(data_poisoning))
      print(data_poisoning.head(10))
```

POISONED DATA. size: 2229

	x	y
3361	No messages on her phone. I'm holding it now	spam
622	MAKE SURE ALEX KNOWS HIS BIRTHDAY IS OVER IN F...	spam
203	Goodmorning sleeping ga.	spam
1889	I gotta collect da car at 6 lei.	spam
423	URGENT! Your Mobile number has been awarded wi...	ham
2493	they released another Italian one today and it...	spam
999	Aight will do, thanks again for comin out	spam
2814	Some are lasting as much as 2 hours. You might...	spam
5048	Edison has rightly said, \A fool can ask more ...	spam
1490	Cant believe i said so many things to you this...	spam

```
[19]: features = cv.fit_transform(pd.concat([x_train,x_poisoning]))
```

```
[21]: model = svm.SVC()  
model.fit(features,pd.concat([y_train,y_poisoning]))
```


```
[21]: SVC()
```

1.7 Evaluate Poisoned Model Accuracy

```
[22]: features_test = cv.transform(x_test)  
print("Accuracy: {}".format(model.score(features_test,y_test)))
```

Accuracy: 0.5273542600896861

```
[23]: y_pred_test = model.predict(features_test)  
  
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)  
plt.show()
```



1.8 SPAM only poisoning

```
[40]: y_spam_only = [ row['v1'] for index, row in spam.iterrows() if row['v1'] ==  
    ↪ 'spam']  
x_spam_only = [ row['v2'] for index, row in spam.iterrows() if row['v1'] ==  
    ↪ 'spam']
```

```
[41]: data_spam_only_dic = {'x': x_spam_only, 'y': y_spam_only}  
data_spam_only = pd.DataFrame(data_spam_only_dic)  
y_spam_only = data_spam_only.y  
x_spam_only = data_spam_only.x  
print("POISONED DATA. size:", len(data_spam_only))  
print(data_spam_only.head(10))
```

POISONED DATA. size: 747

	x	y
0	Free entry in 2 a wkly comp to win FA Cup fina...	spam
1	FreeMsg Hey there darling it's been 3 week's n...	spam
2	WINNER!! As a valued network customer you have...	spam
3	Had your mobile 11 months or more? U R entitle...	spam
4	SIX chances to win CASH! From 100 to 20,000 po...	spam
5	URGENT! You have won a 1 week FREE membership ...	spam
6	XXXMobileMovieClub: To use your credit, click ...	spam
7	England v Macedonia - dont miss the goals/team...	spam
8	Thanks for your subscription to Ringtone UK yo...	spam
9	07732584351 - Rodger Burns - MSG = We tried to...	spam

```
[42]: for i, y in enumerate(y_spam_only):  
    #if random.random() < 0.5:  
    y_spam_only.loc[i] = 'ham'
```

```
[43]: print("POISONED DATA. size:", len(data_spam_only))  
print(data_spam_only.head(10))
```

POISONED DATA. size: 747

	x	y
0	Free entry in 2 a wkly comp to win FA Cup fina...	ham
1	FreeMsg Hey there darling it's been 3 week's n...	ham
2	WINNER!! As a valued network customer you have...	ham
3	Had your mobile 11 months or more? U R entitle...	ham
4	SIX chances to win CASH! From 100 to 20,000 po...	ham
5	URGENT! You have won a 1 week FREE membership ...	ham
6	XXXMobileMovieClub: To use your credit, click ...	ham
7	England v Macedonia - dont miss the goals/team...	ham
8	Thanks for your subscription to Ringtone UK yo...	ham
9	07732584351 - Rodger Burns - MSG = We tried to...	ham

```
[44]: features_spam_only = cv.fit_transform(pd.concat([x_spam_only, x_train]))
```

```
[45]: model_poisoning_spam = svm.SVC()  
model_poisoning_spam.fit(features_spam_only,pd.concat([y_spam_only, y_train]))
```

```
[45]: SVC()
```

1.9 Evaluate Poisoned Model Accuracy

```
[46]: features_test = cv.transform(x_test)  
print("Accuracy: {}".format(model_poisoning_spam.score(features_test,y_test)))
```

Accuracy: 0.8717488789237668

```
[47]: y_pred_test = model_poisoning_spam.predict(features_test)  
  
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)  
plt.show()
```

